

ORACLE  
**OPEN**  
WORLD

#MySQL #ooow16

# How to Analyze and Tune MySQL Queries for Better Performance

Øystein Grøvlen  
Senior Principal Software Engineer  
MySQL Optimizer Team, Oracle

ORACLE®

Copyright © 2016, Oracle and/or its affiliates. All rights reserved.





# Program Agenda

- 1 Cost-based query optimization in MySQL
- 2 Tools for monitoring, analyzing, and tuning queries
- 3 Data access and index selection
- 4 Join optimizer
- 5 Subqueries
- 6 Sorting
- 7 Influencing the optimizer



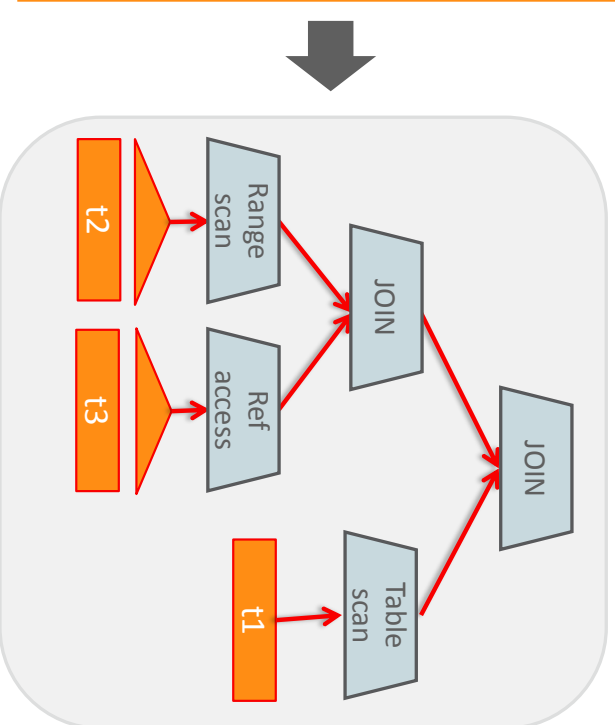
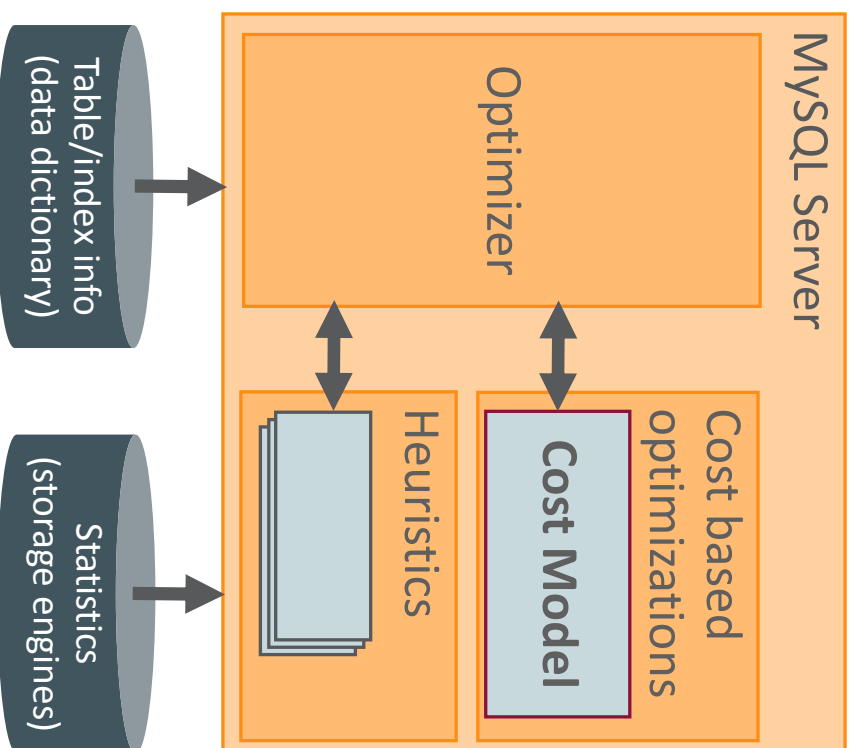
# Program Agenda

- 1 Cost-based query optimization in MySQL
- 2 Tools for monitoring, analyzing, and tuning queries
- 3 Data access and index selection
- 4 Join optimizer
- 5 Subqueries
- 6 Sorting
- 7 Influencing the optimizer

# MySQL Optimizer



```
SELECT a, b
FROM t1, t2, t3
WHERE t1.a = t2.b
AND t2.b = t3.c
AND t2.d > 20
AND t2.d < 30;
```



# Cost-based Query Optimization

## General idea


- Assign cost to operations
- Assign cost to partial or alternative plans
- Search for plan with lowest cost

- Cost-based optimizations:

Access method


Join order

Subquery strategy

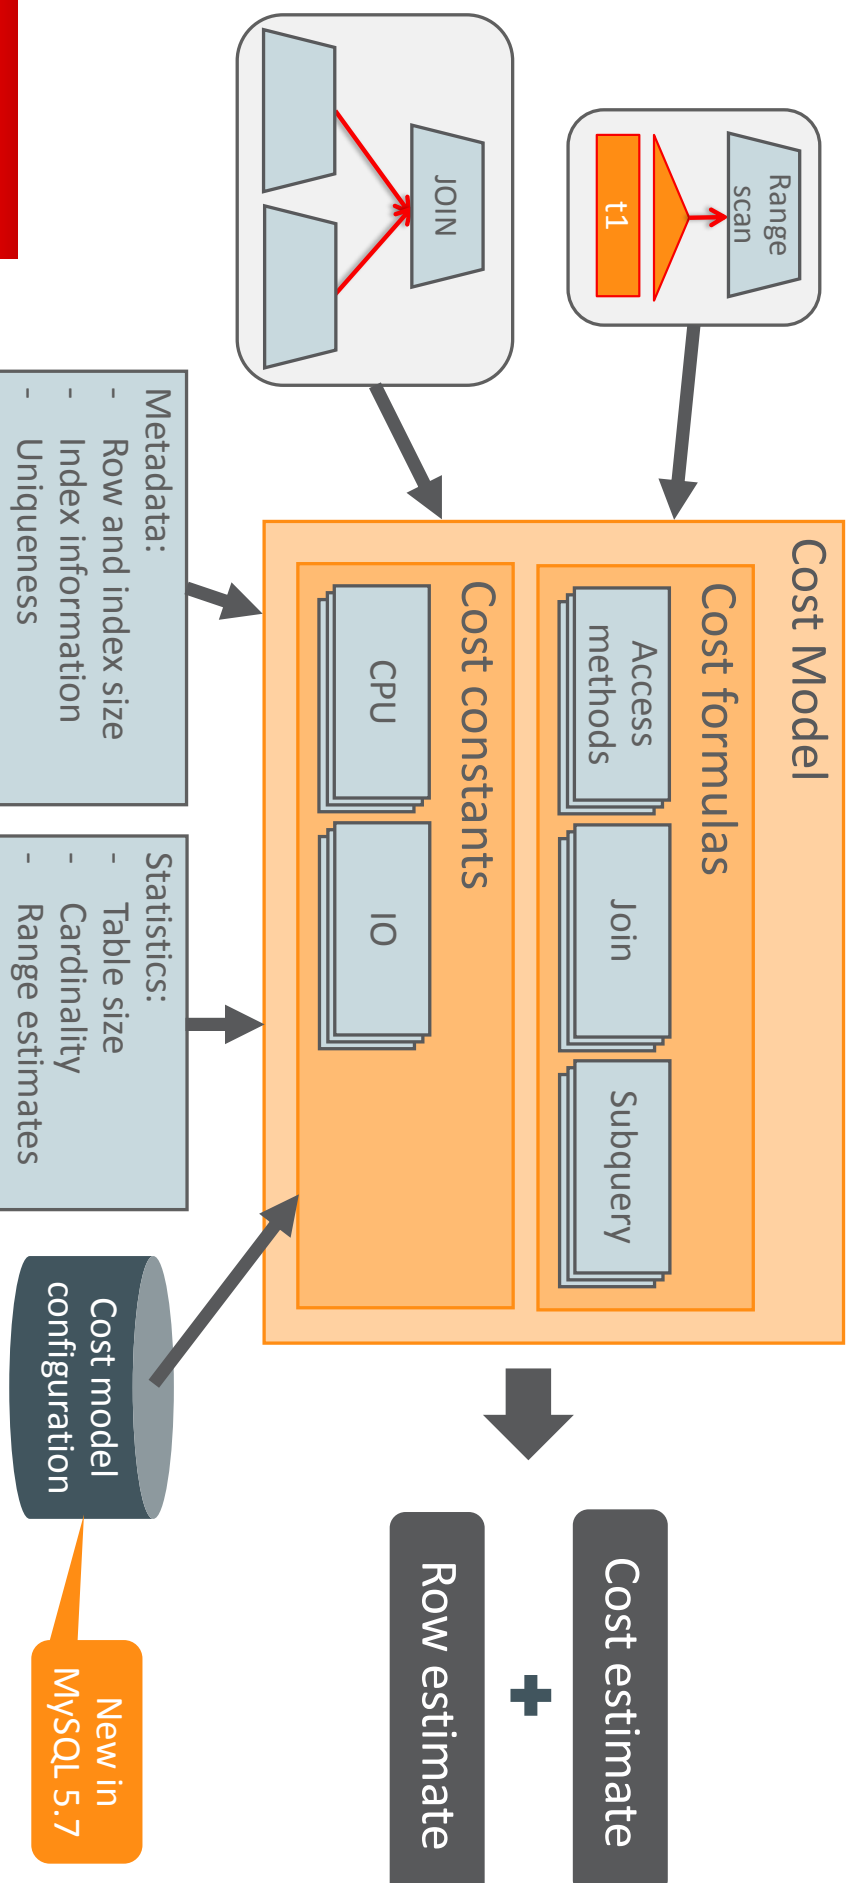


# MySQL Optimizer Characteristics

- Produce the query plan that uses least resources
  - IO and CPU
- Optimizes a single query
  - No inter-query optimizations
- Produces left-deep linear query execution plan



# Optimizer Cost Model



# Cost Estimates

- The **cost** for **executing** a query
- **Cost unit:**
  - “read a random data page from disk”
- Main cost factors:
  - IO cost:
    - #pages read from table
    - #pages read from index
  - CPU cost:
    - Evaluating query conditions
    - Comparing keys/records
    - Sorting keys

- Main cost constants:

Cost	Default value
Reading a random disk page	1.0
Reading a data page from memory buffer	1.0
Evaluating query condition	0.2
Comparing key/record	0.1

MySQL 5.7:  
Configurable




## Cost Model Example

```
SELECT SUM(o_totalprice) FROM orders  
WHERE o_orderdate BETWEEN '1994-01-01' AND '1994-12-31';
```


### Table scan:

- IO-cost: #pages in table \* IO\_BLOCK\_READ\_COST
- CPU cost: #rows \* ROW\_EVALUATE\_COST



### Range scan (on secondary index):

- IO-cost: #rows\_in\_range \* IO\_BLOCK\_READ\_COST
- CPU cost: #rows\_in\_range \* ROW\_EVALUATE\_COST



# Cost Model Example

EXPLAIN SELECT SUM(o\_totalprice) FROM orders

WHERE o\_orderdate BETWEEN '1994-01-01' AND '1994-12-31';

id	select type	table	type	possible keys	key	key len	ref	rows	extra
1	SIMPLE	orders	ALL	i_o_orderdate	NULL	NULL	NULL	15000000	Using where

EXPLAIN SELECT SUM(o\_totalprice) FROM orders

WHERE o\_orderdate BETWEEN '1994-01-01' AND '1994-06-30';

id	select type	table	type	possible keys	key	key len	ref	rows	extra
1	SIMPLE	orders	range	i_o_orderdate	i_o_orderdate	4	NULL	2235118	Using index condition

# Cost Model Example: Optimizer Trace

join\_optimization / row\_estimation / table : orders / range\_analysis

```
"table_scan": {
  "rows": 15000000,
  "cost": 3.12e6
} /* table_scan */,
"potential_range_indices": [
  {
    "index": "PRIMARY",
    "usable": false,
    "cause": "not_applicable"
  },
  {
    "index": "i_o_orderdate",
    "usable": true,
    "key_parts": [ "o_orderDATE", "o_orderkey" ]
  }
] /* potential_range_indices */,
...

"analyzing_range_alternatives": {
  "range_scan_alternatives": [
    {
      "index": "i_o_orderdate",
      "ranges": [ "1994-01-01 <= o_orderDATE <= 1994-12-31"
      ],
      "index_dives_for_eq_ranges": true,
      "rowid_ordered": false,
      "using_mrr": false,
      "index_only": false,
      "rows": 4489990,
      "cost": 5.39e6,
      "chosen": false,
      "cause": "cost"
    }
  ] /* range_scan_alternatives */,
  ...
} /* analyzing_range_alternatives */
```



# Cost Model vs Real World

## Measured Execution Times

	Data in Memory	Data on Disk	Data on SSD
Table scan	6.8 seconds	36 seconds	15 seconds
Index scan	5.2 seconds	2.5 hours	30 minutes

### Force Index Scan:

```
SELECT SUM(o_totalprice)
FROM orders FORCE INDEX (i_o_orderdate)
WHERE o_orderdate BETWEEN '1994-01-01' AND '1994-12-31';
```

# Performance Schema

## Disk I/O

```
SELECT event_name, count_read, avg_timer_read/1000000000.0 "Avg Read Time (ms)",  
       sum_number_of_bytes_read "Bytes Read"  
FROM performance_schema.file_summary_by_event_name  
WHERE event_name='wait/io/file/innodb/innodb_data_file';
```

## Table Scan

event_name	count_read	Avg Read Time (ms)	Bytes Read
wait/io/file/innodb/innodb_data_file	115769	0.0342	1896759296

## Index Scan

event_name	count_read	Avg Read Time (ms)	Bytes Read
wait/io/file/innodb/innodb_data_file	218853	4.2094	35862167552



# Program Agenda

- 1 Cost-based query optimization in MySQL
- 2 Tools for monitoring, analyzing, and tuning queries
- 3 Data access and index selection
- 4 Join optimizer
- 5 Subqueries
- 6 Sorting
- 7 Influencing the optimizer

# Useful tools

- MySQL Enterprise Monitor (MEM), Query Analyzer
  - Commercial product
- Performance schema, MySQL sys schema
- EXPLAIN
  - Tabular EXPLAIN
  - Structured EXPLAIN (FORMAT=JSON)
  - Visual EXPLAIN (MySQL Workbench)
- Optimizer trace
- Slow log
- Status variables (SHOW STATUS LIKE 'Sort%')

# MySQL Enterprise Monitor, Query Analyzer

ORACLE MySQL Enterprise Monitor

Dashboards ▾ Events Query Analyzer Reports & Graphs ▾ Configuration ▾

1 2 0 0 1 admin ▾ 1 2

Refresh: Every Minute ▾

Browse Queries

Show 25 entities Export data options...

Showing 1 to 25 of 1,197 entries First Previous 1 2 3 4 5 Next Last

Query	Database	Counts			QRTI	Latency (h:mm:ss.ms)					Rows		
		Exec	Err	Warn		Total	Max	Avg	Locks	Avg History	Total	Examined	
COMMIT (1)	mem	24,707	0	0	0.92	20:27.872	5.828	0.050	0.000		0	0	
INSERT INTO `mem`...o`), `hostTo`), ... (1)	mem	6,903	0	0	0.79	18:16.538	17.784	0.159	6.699		7,356	0	
INSERT INTO `mem`...on` = IF (VALUES ( ... (1)	mem	6,985	0	0	0.86	10:54.856	8.940	0.094	8.301		7,332	0	
INSERT INTO `mem`...uan`...en`), `lastSeen` ) (1)	mem	7,025	37	0	1.00	3:11.791	11.220	0.027	4.436		13,947	0	
INSERT INTO `mem`...instr...teny`), `latency` ) (1)	mem	740	0	0	0.77	1:46.459	8.745	0.147	0.147		1,386	0	
SELECT `mysqldconn0`...asProc18, 1191, 0, ... (1)	mem	974	0	0	0.90	1:01.829	12.359	0.063	0.353		974	974	
SELECT `mysqldserve0`...s`, `hostCache` AS ... (1)	mem	1,801	0	0	0.93	53.661	13.351	0.030	0.636		1,801	1,801	
INSERT INTO `mem`...instr...es`), `diskWrites` ) (1)	mem	26	0	0	0.79	44.838	14.954	1.725	0.070		26	0	
UPDATE `mem`...inventory`...n` = ? WHERE `hid` = ? (1)	mem	321	0	0	0.80	43.886	9.668	0.137	0.044		321	321	
UPDATE `mem`...config`...? WHERE `user_id` = ? (1)	mem	321	0	0	0.71	40.788	10.607	0.127	0.038		321	321	
CREATE TEMPORARY TABLE ...(`id` INT8 NOT NULL) (1)	mem	13	0	0	0.14	36.525	12.055	2.810	0.000		0	0	
INSERT INTO `mem`...instr...nedTableDefinitions` ) (1)	mem	26	0	0	0.64	34.348	13.349	1.321	0.003		27	0	
UPDATE `mem`...inventory`...p` = ? WHERE `hid` = ? (1)	mem	416	0	0	0.81	33.469	13.060	0.080	0.042		416	416	
INSERT INTO `mem`...instr...hed`), `notCached` ) (1)	mem	26	0	0	0.61	32.509	11.782	1.250	0.003		27	0	
SELECT * FROM (SELECT ...m_no_index_used` AS ... (1)	mem	14	0	0	0.00	29.832	9.602	2.131	0.008		7,459	227,050	
INSERT INTO `mem`...instr...s`), `connections` ) (1)	mem	25	0	0	0.82	29.462	14.294	1.178	0.005		26	0	
INSERT INTO `mem`...instr...(`sent`), `sent` ) (1)	mem	25	0	0	0.82	28.991	14.332	1.160	0.003		26	0	
SELECT `agent0`...`hid`...`hid` = ? FOR UPDATE (1)	mem	909	0	0	0.96	28.632	6.964	0.031	0.158		909	909	
UPDATE `mem`...events`...me` = ? WHERE `id` = ? (1)	mem	395	0	0	0.92	27.544	10.061	0.070	0.038		395	395	

Copyright © 2005, 2013, Oracle and/or its affiliates. All rights reserved.

3.0.1.7150 - Cerberus.local (192.168.1.67) - Sep 16, 2013 11:50:28 am (Up Since: 38 minutes ago) - About

ORACLE

Copyright © 2016, Oracle and/or its affiliates. All rights reserved. |





# Query Analyzer Query Details

Canonical Query   **Example Query**   Explain Query   Graphs

The query with the longest execution time during the Time Span (usually the slowest but not always).

## Sampled Query

truncated | full | formatted

```
SELECT
  mysqlserv0.`hid` AS hid1124_0, mysqlserv0.`id` AS id2_1124_0,
  mysqlserv0.`lastContact` AS lastContact3_1124_0,
  mysqlserv0.`lastContact` AS lastContact4_1124_0,
  mysqlserv0.`startTime` AS startTime5_1124_0,
  mysqlserv0.`hasStartTime` AS hasStart6_1124_0,
  mysqlserv0.`timestamp` AS timestamp7_1124_0,
  mysqlserv0.`capabilities` AS capabil18_1124_0,
  mysqlserv0.`capabilities` AS capab118_1124_0,
  mysqlserv0.`characterSet` AS charact10_1124_0,
  mysqlserv0.`characterSet` AS hasChar11_1124_0,
  mysqlserv0.`hasCharacterSet` AS hasColl12_1124_0,
  mysqlserv0.`collation` AS hasColl13_1124_0,
  mysqlserv0.`connection` AS connection14_1124_0,
  mysqlserv0.`hasConnection` AS hasConn15_1124_0,
  mysqlserv0.`environment` AS environ16_1124_0,
```

**Execution Time**  
27,084 ms

**Date**  
Sep 16, 2013 1:07:17 PM

**User**  
service\_manager

**Thread ID**  
10,712

**From Host**  
localhost




**To Host**  
localhost

**Source Location**  
None found.

**Comments**  
None found.

Canonical Query   Example Query   Explain Query   **Graphs**

Graphs of a query during the Time Span



# Performance Schema

## Some useful tables

- `events_statements_history`  
`events_statements_history_long`
  - Most recent statements executed
- `events_statements_summary_by_digest`
  - Summary for similar statements (same statement digest)
- `file_summary_by_event_name`
  - Interesting event: `wait/io/file/innodb/innodb_data_file`
- `table_io_waits_summary_by_table`  
`table_io_waits_summary_by_index_usage`
  - Statistics on storage engine access per table and index

# Performance Schema

## Statement events

- Tables:
  - `events_statements_current` (Current statement for each thread)
  - `events_statements_history` (10 most recent statements per thread)
  - `events_statements_history_long` (10000 most recent statements)

- Columns:

`THREAD_ID`, `EVENT_ID`, `END_EVENT_ID`, `EVENT_NAME`, `SOURCE`, `TIMER_START`, `TIMER_END`, `TIMER_WAIT`, `LOCK_TIME`, `SQL_TEXT`, `DIGEST`, `DIGEST_TEXT`, `CURRENT_SCHEMA`, `OBJECT_TYPE`, `OBJECT_SCHEMA`, `OBJECT_NAME`, `OBJECT_INSTANCE_BEGIN`, `MYSQL_ERRNO`, `RETURNED_SQLSTATE`, `MESSAGE_TEXT`, `ERRORS`, `WARNINGS`, `ROWS_AFFECTED`, `ROWS_SENT`, `ROWS_EXAMINED`, `CREATED_TMP_DISK_TABLES`, `CREATED_TMP_TABLES`, `SELECT_FULL_JOIN`, `SELECT_FULL_RANGE_JOIN`, `SELECT_RANGE_CHECK`, `SELECT_SCAN`, `SORT_MERGE_PASSES`, `SORT_RANGE`, `SORT_ROWS`, `SORT_SCAN`, `NO_INDEX_USED`, `NO_GOOD_INDEX_USED`, `NESTING_EVENT_ID`, `NESTING_EVENT_TYPE`

# Performance Schema

## Statement digest

- Normalization of queries to group statements that are similar to be grouped and summarized:

```
SELECT * FROM orders WHERE o_custkey=10 AND o_totalprice>20  
SELECT * FROM orders WHERE o_custkey = 20 AND o_totalprice > 100
```

➡ **SELECT \* FROM orders WHERE o\_custkey = ? AND o\_totalprice > ?**

- **events\_statements\_summary\_by\_digest**

```
DIGEST, DIGEST_TEXT, COUNT_STAR, SUM_TIMER_WAIT, MIN_TIMER_WAIT, AVG_TIMER_WAIT,  
MAX_TIMER_WAIT, SUM_LOCK_TIME, SUM_ERRORS, SUM_WARNINGS, SUM_ROWS_AFFECTED,  
SUM_ROWS_SENT, SUM_ROWS_EXAMINED, SUM_CREATED_TMP_DISK_TABLES, SUM_CREATED_TMP_TABLES,  
SUM_SELECT_FULL_JOIN, SUM_SELECT_FULL_RANGE_JOIN, SUM_SELECT_RANGE, SUM_SELECT_RANGE_CHECK,  
SUM_SELECT_SCAN, SUM_SORT_MERGE_PASSES, SUM_SORT_RANGE, SUM_SORT_ROWS, SUM_SORT_SCAN,  
SUM_NO_INDEX_USED, SUM_NO_GOOD_INDEX_USED, FIRST_SEEN, LAST_SEEN
```



# MySQL sys Schema

- A collection of views, procedures and functions, designed to make reading raw Performance Schema data easier
- Implements many common DBA and Developer use cases
  - File IO usage per user
  - Which indexes is never used?
  - Which queries use full table scans?
- Examples of very useful functions:
  - `format_time()`, `format_bytes()`, `format_statement()`
- **Included with MySQL 5.7**
- Bundled with MySQL Workbench



# MySQL sys Schema

## Example

**statement\_analysis:** Lists a normalized statement view with aggregated statistics, ordered by the total execution time per normalized statement

```
mysql> SELECT * FROM sys.statement_analysis LIMIT 1\G
*****
query: INSERT INTO `mem__quan`.``nor ... nDuration` = IF ( VALUES ( ...
db: mem
full_scan: 0
exec_count: 1110067
err_count: 0
warn_count: 0
total_latency: 1.93h
max_latency: 5.03 s
avg_latency: 6.27 ms
lock_latency: 00:18:29.18

rows_sent: 0
rows_sent_avg: 0
rows_examined: 0
rows_examined_avg: 0
tmp_tables: 0
tmp_disk_tables: 0
rows_sorted: 0
sort_merge_passes: 0
digest: d48316a218e95b1b8b72db5e6b177788!
first_seen: 2014-05-20 10:42:17
```

# EXPLAIN

Understand the query plan

- Use **EXPLAIN** to print the final query plan:

```
EXPLAIN SELECT * FROM t1 JOIN t2 ON t1.a = t2.a WHERE b > 10 AND c > 10;
```

id	select_ type	table	partitions	type	possible keys	key	key len	ref	rows	filtered	Extra
1	SIMPLE	t1	NULL	range	PRIMARY, idx1	idx1	4	NULL	12	33.33	Using index condition
2	SIMPLE	t2	NULL	ref	idx2	idx2	4	t1.a	1	100.00	NULL

- Explain for a running query (**New in MySQL 5.7**):  
**EXPLAIN FOR CONNECTION *connection\_id***;

# Structured EXPLAIN

- JSON format:

**EXPLAIN FORMAT=JSON SELECT ...**

- Contains more information:

- Used index parts
- Pushed index conditions
- Cost estimates
- Data estimates

Added in  
MySQL 5.7

```
EXPLAIN FORMAT=JSON
SELECT * FROM t1 WHERE b > 10 AND c > 10;
EXPLAIN
{
  "query_block": {
    "select_id": 1,
    "cost_info": {
      "query_cost": "17.81"
    },
    "table": {
      "table_name": "t1",
      "access_type": "range",
      "possible_keys": [
        "idx1"
      ],
      "key": "idx1",
      "used_key_parts": [
        "b"
      ],
      "key_length": "4",
      "rows_examined_per_scan": 12,
      "rows_produced_per_join": 3,
      "filtered": "33.33%",
      "index_condition": "('test`.`t1`.`b` > 10)",
      "cost_info": {
        "read_cost": "17.01",
        "eval_cost": "0.80",
        "prefix_cost": "17.81",
        "data_read_per_join": "63"
      },
      "attached_condition": "('test`.`t1`.`c` > 10)"
    }
  }
}
```



# Structured EXPLAIN

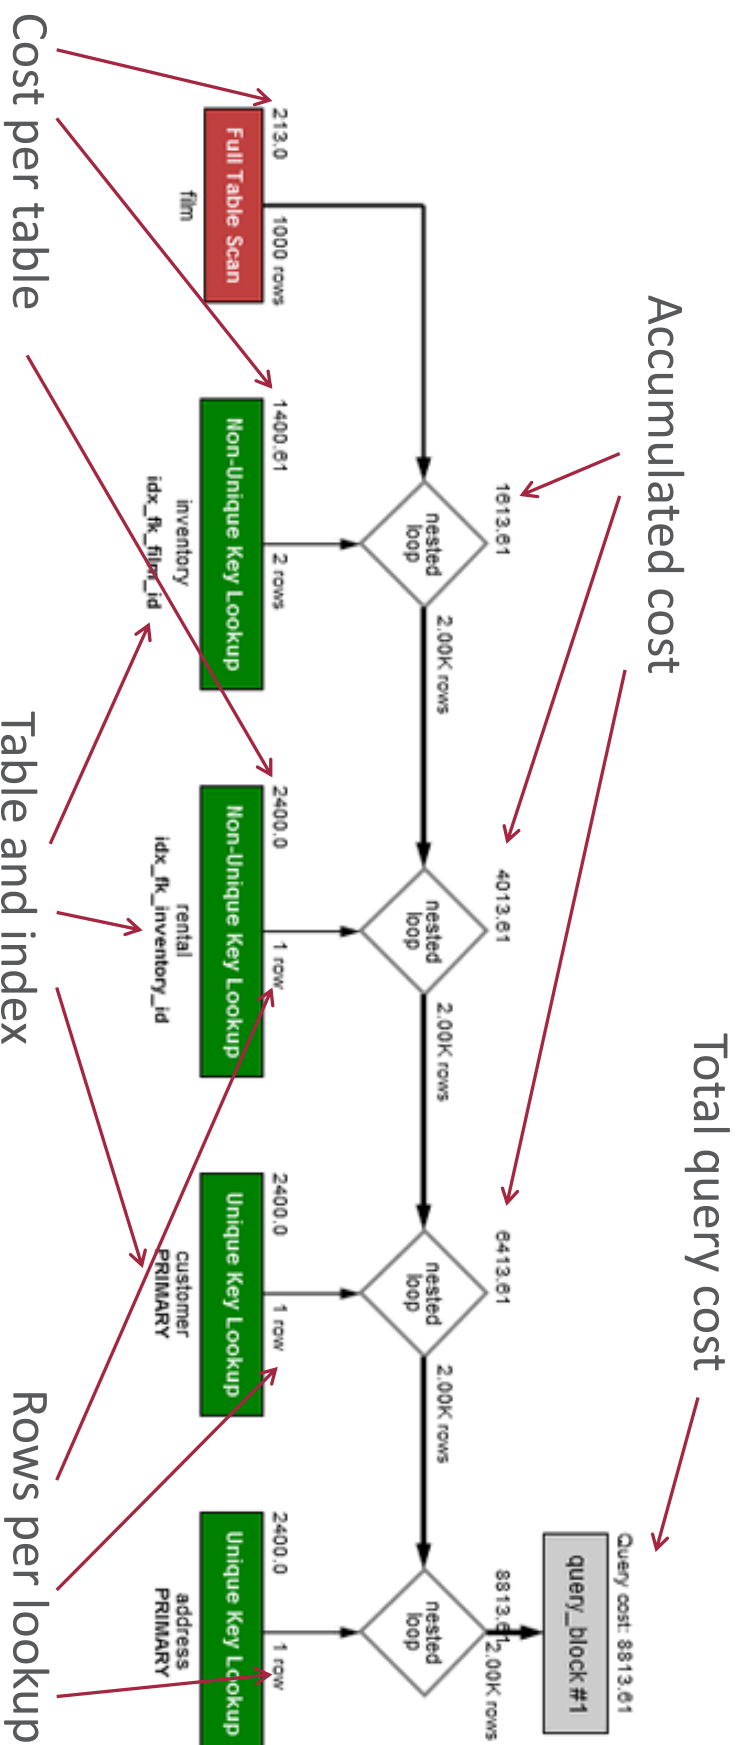
## Assigning Conditions to Tables

EXPLAIN **FORMAT=JSON** SELECT \* FROM t1, t2  
WHERE t1.a=t2.a AND t2.a=9 AND (NOT (t1.a > 10 OR t2.b >3) OR (t1.b=t2.b+7 AND t2.b = 5));

### EXPLAIN

```
{
  "query_block": {
    "select_id": 1,
    "nested_loop": [
      {
        "table": {
          "table_name": "t1",
          "access_type": "ALL",
          "rows": 10,
          "filtered": 100,
          "attached_condition": "(t1.a = 9)"
        } /* table */
      },
      {
        "table": {
          "table_name": "t2",
          "access_type": "ALL",
          "rows": 10,
          "filtered": 100,
          "using_join_buffer": "Block Nested Loop",
          "attached_condition": "(((t2.a = 9) and ((t2.b <= 3) or ((t2.b = 5) and (t1.b = 12)))))"
        } /* table */
      } /* nested_loop */
    ] /* query_block */
  }
}
```

# Visual EXPLAIN



# Optimizer Trace: Query Plan Debugging

- EXPLAIN shows the selected plan
- Optimizer trace shows WHY the plan was selected

```
SET optimizer_trace="enabled=on";
```

```
SELECT * FROM t1,t2 WHERE f1=1 AND f1=f2 AND f2>0;
```

```
SELECT trace FROM information_schema.optimizer_trace  
INTO OUTFILE '<filename>' LINES TERMINATED BY ";
```

```
SET optimizer_trace="enabled=off";
```

QUERY	SELECT * FROM t1,t2 WHERE f1=1 AND f1=f2 AND f2>0;
TRACE	"steps": [ { "join_preparation": { "select#": 1,... } ... } ... ]
MISSING_BYTES_BEYOND_MAX_MEM_SIZE	0
INSUFFICIENT_PRIVILEGES	0

# Optimizer Trace

join\_optimization / row\_estimation / table : orders / range\_analysis

```
"table_scan": {
  "rows": 15000000,
  "cost": 3.12e6
} /* table_scan */

"potential_range_indices": [
  {
    "index": "PRIMARY",
    "usable": false,
    "cause": "not_applicable"
  },
  {
    "index": "i_o_orderdate",
    "usable": true,
    "key_parts": [ "o_orderDATE", "o_orderkey" ]
  }
] /* potential_range_indices */
...

"analyzing_range_alternatives": {
  "range_scan_alternatives": [
    {
      "index": "i_o_orderdate",
      "ranges": [ "1994-01-01 <= o_orderDATE <= 1994-12-31"
      ],
      "index_dives_for_eq_ranges": true,
      "rowid_ordered": false,
      "using_mrr": false,
      "index_only": false,
      "rows": 4489990,
      "cost": 5.39e6,
      "chosen": false,
      "cause": "cost"
    }
  ] /* range_scan_alternatives */
  ...
} /* analyzing_range_alternatives */
```



# Program Agenda

- 1 Cost-based query optimization in MySQL
- 2 Tools for monitoring, analyzing, and tuning queries
- 3 Data access and index selection**
- 4 Join optimizer
- 5 Subqueries
- 6 Sorting
- 7 Influencing the optimizer

# Selecting Access Method

**Finding the optimal method to read data from storage engine**

- For each table, find the best access method:
  - Check if the access method is useful
  - Estimate cost of using access method
  - Select the cheapest to be used
- Choice of access method is cost based

## Main access methods:

- Table scan
- Index scan
- Index look-up (ref access)
- Range scan
- Index merge
- Loose index scan

# Ref Access

## Single Table Queries

**EXPLAIN SELECT \* FROM customer WHERE c\_custkey = 570887;**

id	select type	table	type	possible keys	key	key len	ref	rows	filtered	Extra
1	SIMPLE	customer	const	PRIMARY	PRIMARY	4	const	1	100.00	NULL

**EXPLAIN SELECT \* FROM orders WHERE o\_orderdate = '1992-09-12';**

id	select type	table	type	possible keys	key	key len	ref	rows	filtered	Extra
1	SIMPLE	orders	ref	i_o_orderdate	i_o_orderdate	4	const	6272	100.00	NULL

# Ref Access

## Join Queries

```
EXPLAIN SELECT *
FROM orders JOIN customer ON c_custkey = o_custkey
WHERE o_orderdate = '1992-09-12';
```

id	select type	table	type	possible keys	key	key len	ref	rows	filtered	Extra
1	SIMPLE	orders	ref	i_o_orderdate, i_o_custkey	i_o_orderdate	4	const	6272	100.00	Using where
1	SIMPLE	customer	eq_ref	PRIMARY	PRIMARY	4	dbt3. orders. o_custkey	1	100.00	NULL



# Ref Access

Join Queries, continued

```
EXPLAIN SELECT *
FROM orders JOIN customer ON c_custkey = o_custkey
WHERE c_acctbal < -1000;
```

id	select type	table	type	possible keys	key	key len	ref	rows	filtered	Extra
1	SIMPLE	customer	ALL	PRIMARY	NULL	NULL	NULL	1500000	33.33	Using where
1	SIMPLE	orders	ref	i_o_custkey	i_o_custkey	5	dbt3. customer. c_custkey	7	100.00	NULL

# Range Optimizer

- Goal: find the "minimal" ranges for each index that needs to be read
- Example:  
**`SELECT * FROM t1 WHERE (key1 > 10 AND key1 < 20) AND key2 > 30`**
- Range scan using INDEX(key1):



- Range scan using INDEX(key2):



## Range Optimizer, cont.

- Range optimizer selects the “useful” parts of the WHERE condition:
  - Conditions comparing a column value with a constant:

<b>key &gt; 3</b>	<b>key BETWEEN 4 AND 6</b>	<b>key IS NULL</b>
<b>key = 4</b>	<b>key IN (10,12,..)</b>	<b>key LIKE "abc%"</b>

- Nested AND/OR conditions are supported

- Result: list of disjoint ranges that need to be read from index:



- Cost estimate based on number of records in each range:
  - Record estimate is found by asking the Storage Engine (“index dives”)



# Range Optimizer

Optimizer Trace show ranges

```
SELECT a, b FROM t1
WHERE a > 10
AND a < 25
AND a NOT IN (11, 19))
AND (b < 5 OR b > 10);
```

```
"analyzing_range_alternatives": {
  "range_scan_alternatives": [
    {
      "index": "i_a",
      "ranges": [
        "10 < a < 11",
        "11 < a < 19",
        "19 < a < 25"
      ],
      "index_dives_for_eq_ranges": true,
      "rowid_ordered": false,
      "using_mrr": false,
      "index_only": false,
      "rows": 3,
      "cost": 6.61,
      "chosen": true
    },
    {
      "index": "i_b",
      "ranges": [
        "NULL < b < 5",
        "10 < b"
      ],
      "index_dives_for_eq_ranges": true,
      "rowid_ordered": false,
      ...
    }
  ]
}
```

# Range Optimizer: Case Study

Why table scan?

```
SELECT * FROM orders
WHERE YEAR(o_orderdate) = 1997 AND MONTH(o_orderdate) = 5
AND o_clerk = 'Clerk#000001866';
```

id	select type	table	type	possible keys	key	key len	ref	rows	Extra
1	SIMPLE	orders	ALL	NULL	NULL	NULL	NULL	15000000	Using where

Index not considered

```
mysql> SELECT * FROM orders WHERE year(o_orderdate) = 1997 AND MONTH(...
...
15 rows in set (8.91 sec)
```

## Some Reasons Why Index can not be Used

- Indexed column is used as argument to function  
`YEAR(o_orderdate) = 1997`
- Looking for a suffix:  
`name LIKE '%son'`
- First column(s) of compound index NOT used  
`b = 10` when index defined over `(a, b)`
- Type mismatch  
`my_string = 10`
- Character set / collation mismatch  
`t1 LEFT JOIN t2 ON t1.utf8_string = t2.latin1_string`



# Range Optimizer: Case Study

Rewrite query to avoid functions on indexed columns

```
SELECT * FROM orders
WHERE o_orderdate BETWEEN '1997-05-01' AND '1997-05-31'
AND o_clerk = 'Clerk#000001866';
```

id	select type	table	type	possible keys	key	key len	ref	rows	Extra
1	SIMPLE	orders	range	i_o_orderdate	i_o_orderdate	4	NULL	376352	Using index condition; Using where

```
mysql> SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' AND ...
...
15 rows in set (0.91 sec)
```

ORACLE®



# Range Optimizer: Case Study

Adding another index

**CREATE INDEX i\_o\_clerk ON orders(o\_clerk);**

```
SELECT * FROM orders
WHERE o_orderdate BETWEEN '1997-05-01' AND '1997-05-31'
AND o_clerk = 'Clerk#000001866';
```

id	select type	table	type	possible keys	key	key len	ref	rows	Extra
1	SIMPLE	orders	range	i_o_orderdate, i_o_clerk	i_o_clerk	16	NULL	1504	Using index condition; Using where

```
mysql> SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' AND ...
...
15 rows in set (0.01 sec)
```

ORACLE®



# Range Access for Multi-Column Indexes

Example table with multi-part index

- Table:


pk	a	b	c
----	---	---	---



- INDEX idx(a, b, c);



- Logical storage layout of index:




## Range Access for Multi-Column Indexes, cont

- Equality on 1<sup>st</sup> index column?
  - Can add condition on 2<sup>nd</sup> index column to range condition

- Example:

**SELECT \* from t1 WHERE a IN (10,11,13) AND (b=2 OR b=4)**




- Resulting range scan:



# Range Access for Multi-Column Indexes, cont

- Non-Equality on 1<sup>st</sup> index column:
  - Can **NOT** add condition on 2<sup>nd</sup> index column to range condition
- Example:

**SELECT \* from t1 WHERE a > 10 AND a < 13 AND (b=2 OR b=4)**



- Resulting range scan:



# Range Optimizer: Case Study

Create multi-column index

**CREATE INDEX i\_o\_clerk\_date ON orders(o\_clerk, o\_orderdate);**

```
SELECT * FROM orders
WHERE o_orderdate BETWEEN '1997-05-01' AND '1997-05-31'
AND o_clerk = 'Clerk#000001866';
```

id	select type	table	type	possible keys	key	key len	ref	rows	Extra
1	SIMPLE	orders	range	i_o_orderdate, i_o_clerk, i_o_clerk_date	i_o_clerk_date	20	NULL	14	Using index condition

```
mysql> SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' AND ...
...
15 rows in set (0.00 sec)
```



# Performance Schema: Query History

**UPDATE performance\_schema.setup\_consumers**  
**SET enabled='YES' WHERE name = 'events\_statements\_history';**

MySQL 5.7:  
Enabled by default


```
mysql> SELECT sql_text, (timer_wait)/1000000000.0 "t (ms)", rows_examined rows
        FROM performance_schema.events_statements_history ORDER BY timer_start;
+-----+-----+-----+-----+-----+-----+
| sql_text | t (ms) | rows |
+-----+-----+-----+
| SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ... | 8.1690 | 1505 |
| SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ... | 7.2120 | 1505 |
| SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ... | 8.1613 | 1505 |
| SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ... | 7.0535 | 1505 |
| CREATE INDEX i_o_clerk_date ON orders(o_clerk,o_orderdate) | 82036.4190 | 0 |
| SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ... | 0.7259 | 15 |
| SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ... | 0.5791 | 15 |
| SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ... | 0.5423 | 15 |
| SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ... | 0.6031 | 15 |
| SELECT * FROM orders WHERE o_orderdate BETWEEN '1997-05-01' ... | 0.2710 | 15 |
+-----+-----+-----+
```

# Index Merge

- Uses multiple indexes on the same table
- Implemented index merge strategies:
  - **Index Merge Union**
    - OR-ed conditions between different indexes
  - **Index Merge Intersection**
    - AND conditions between different indexes
  - **Index Merge Sort-Union**
    - OR-ed conditions where condition is a range

# Index Merge Union


- Single index cannot handle ORed conditions on different columns
- Example:  
**SELECT \* FROM t1 WHERE a=10 OR b=10**
- Index Merge Union:



# Index Merge Intersection

- Combine several indexes to reduce number of (or avoid) accesses to base table for ANDed conditions
- Example:  
**SELECT \* FROM t1 WHERE a=10 AND b=10**

- Index Merge Intersection:





# Index Merge Intersection: Example 1



```
SELECT COUNT(*) FROM lineitem
WHERE l_shipdate = '1997-05-01' AND l_commitdate = '1997-05-01';
```

id	select type	table	type	possible keys	key	key len	ref	rows	Extra
1	SIMPLE	lineitem	index_merge	i_l_shipdate, i_l_commitdate	i_l_shipdate, i_l_commitdate	4,4	NULL	43	Using intersect (i_l_shipdate, i_l_commitdate); Using where; Using index

```
mysql> SELECT COUNT(*) FROM lineitem WHERE l_shipdate = '1997-05-01' ...
...
1 row in set (0.02 sec)
mysql> SET optimizer_switch='index_merge_intersection=off';
mysql> SELECT COUNT(*) FROM lineitem WHERE l_shipdate = '1997-05-01' ...
...
1 row in set (0.11 sec)
```

# Index Merge Intersection: Example 2

Beware of low-selectivity indexes!

Low selectivity

SELECT count(\*) FROM user  
WHERE user\_type=2 AND status=1 AND parent\_id=0;

id	select type	table	type	possible keys	key	key len	ref	rows	Extra
1	SIMPLE	user	index_merge	parent_id, status, user_type	user_type, status, parent_id	1,1,4	NULL	3696	Using intersect (user_type, status, parent_id); Using where; Using index

```
mysql> SELECT count(*) FROM user WHERE user_type=2 AND status=1 ...
...
1 row in set (5.33 sec)
mysql> SELECT count(*) FROM user USE INDEX (user_type) WHERE user_type=2 ...
...
1 row in set (0.09 sec)
```



# Index Merge Intersection: Example 2

## Handler status variables

```
mysql> FLUSH STATUS;
mysql> SELECT count(*) FROM user
WHERE user_type=2 AND status=1
AND parent_id=0;
```

...

```
1 row in set (5.28 sec)
```

```
mysql> SHOW STATUS LIKE 'Handler_read%';
```

Variable_name	Value
Handler_read_first	0
Handler_read_key	3
Handler_read_last	0
Handler_read_next	15825904
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_next	0

```
mysql> FLUSH STATUS;
mysql> SELECT count(*) FROM user
USE INDEX (user_id) WHERE user_type=2 AND
status=1 AND parent_id=0;
```

...

```
1 row in set (0.09 sec)
```

```
mysql> SHOW STATUS LIKE 'Handler_read%';
```

Variable_name	Value
Handler_read_first	0
Handler_read_key	1
Handler_read_last	0
Handler_read_next	23312
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_next	0




# Program Agenda

- 1 Cost-based query optimization in MySQL
- 2 Tools for monitoring, analyzing, and tuning queries
- 3 Data access and index selection
- 4 **Join optimizer**
- 5 Subqueries
- 6 Sorting
- 7 Influencing the optimizer

# Join Optimizer

## “Greedy search strategy”

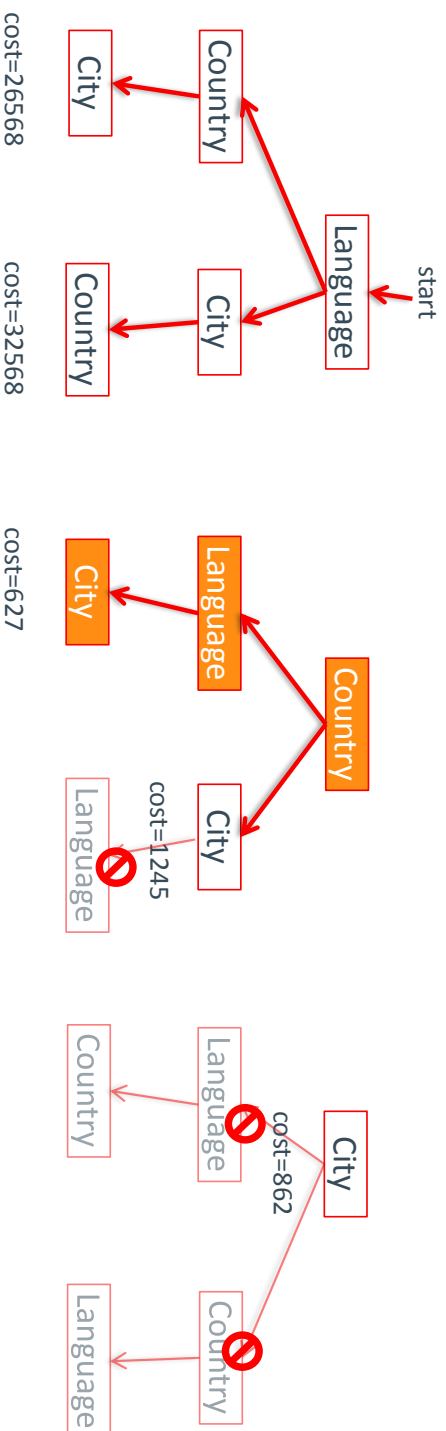
- Goal: Given a JOIN of N tables, find the best JOIN ordering
- Strategy:
  - Start with all 1-table plans (Sorted based on size and *key dependency*)
  - Expand each plan with remaining tables
    - Depth-first
  - If “cost of partial plan” > “cost of best plan”:
    - “prune” plan
  - Heuristic pruning:
    - Prune less promising partial plans
    - May in rare cases miss most optimal plan (turn off with `set optimizer_prune_level = 0`)



N! possible plans

# JOIN Optimizer Illustrated

```
SELECT City.Name, Language FROM Language, Country, City
WHERE City.CountryCode = Country.Code
AND City.ID = Country.Capital
AND City.Population >= 1000000
AND Language.Country = Country.Code;
```





# Join Optimizer

## Example

```
EXPLAIN SELECT *  
FROM customers JOIN orders ON c_custkey = o_custkey  
WHERE c_acctbal < -1000 AND o_orderdate < '1993-01-01';
```

id	select type	table	type	possible keys	key	key len	ref	rows	filtered	Extra
1	SIMPLE	orders	ALL	i_o_orderdate, i_o_custkey	NULL	NULL	NULL	150000000	31.19	Using where
1	SIMPLE	customer	eq_ref	PRIMARY	PRIMARY	4	dbt3.orders. o_custkey	1	33.33	Using where



# Join Optimizer


Change join order with STRAIGHT\_JOIN

```
EXPLAIN SELECT STRAIGHT_JOIN *  
FROM customer JOIN orders ON c_custkey = o_custkey  
WHERE c_acctbal < -1000 AND o_orderdate < '1993-01-01';
```

id	select type	table	type	possible keys	key	key len	ref	rows	filtered	Extra
1	SIMPLE	customer	ALL	PRIMARY	NULL	NULL	NULL	1500000	31.19	Using where
1	SIMPLE	orders	ref	i_o_orderdate, i_o_custkey	i_o_custkey	5	dbt3. customer. c_custkey	15	33.33	Using where



# Join Order Performance



■ orders → customer ■ customer → orders



# Join Order Hints

MySQL 8.0 Optimizer Labs Release

```
EXPLAIN SELECT /*+ JOIN_ORDER(customer, orders) */ *  
FROM customer JOIN orders ON c_custkey = o_custkey  
WHERE c_acctbal < -1000 AND o_orderdate < '1993-01-01';
```

id	select type	table	type	possible keys	key	key len	ref	rows	filtered	Extra
1	SIMPLE	customer	ALL	PRIMARY	NULL	NULL	NULL	1500000	31.19	Using where
1	SIMPLE	orders	ref	i_o_orderdate, i_o_custkey	i_o_custkey	5	dbt3. customer. c_custkey	15	33.33	Using where

Alternatives with same effect for this query:

**JOIN\_PREFIX(customer) JOIN\_SUFFIX(orders) JOIN\_FIXED\_ORDER()**

ORACLE®

## DBT-3 Query 8

### National Market Share Query

```
SELECT o_year, SUM(CASE WHEN nation = 'FRANCE' THEN volume ELSE 0 END) / SUM(volume) AS
mkt_share
FROM (
    SELECT EXTRACT(YEAR FROM o_orderdate) AS o_year,
           l_extendedprice * (1 - l_discount) AS volume, n2.n_name AS nation
    FROM part
    JOIN lineitem ON p_partkey = l_partkey
    JOIN supplier ON s_suppkey = l_suppkey
    JOIN orders ON l_orderkey = o_orderkey
    JOIN customer ON o_custkey = c_custkey
    JOIN nation n1 ON c_nationkey = n1.n_nationkey
    JOIN region ON n1.n_regionkey = r_regionkey
    JOIN nation n2 ON s_nationkey = n2.n_nationkey
    WHERE r_name = 'EUROPE' AND o_orderdate BETWEEN '1995-01-01' AND '1996-12-31'
    AND p_type = 'PROMO BRUSHED STEEL'
) AS all_nations GROUP BY o_year ORDER BY o_year;
```

# DBT-3 Query 8

# MySQL Workbench: Visual EXPLAIN (MySQL 5.6)

Execution time: 21 seconds



## DBT-3 Query 8

Force early processing of high selectivity conditions

```

SELECT o_year, SUM(CASE WHEN nation = 'FRANCE' THEN volume ELSE 0 END) / SUM(volume) AS
mkt_share
FROM (
    SELECT EXTRACT(YEAR FROM o_orderdate) AS o_year,
           l_extendedprice * (1 - l_discount) AS volume, n2.n_name AS nation
    FROM part
    STRAIGHT JOIN lineitem ON p_partkey = l_partkey
    JOIN supplier ON s_suppkey = l_suppkey
    JOIN orders ON l_orderkey = o_orderkey
    JOIN customer ON c_custkey = o_custkey
    JOIN nation n1 ON c_nationkey = n1.n_nationkey
    JOIN region ON n1.n_regionkey = r_regionkey
    JOIN nation n2 ON s_nationkey = n2.n_nationkey
    WHERE r_name = 'EUROPE' AND o_orderdate BETWEEN '1995-01-01' AND '1996-12-31'
    AND p_type = 'PROMO BRUSHED STEEL'
) AS all_nations GROUP BY o_year ORDER BY o_year;

```

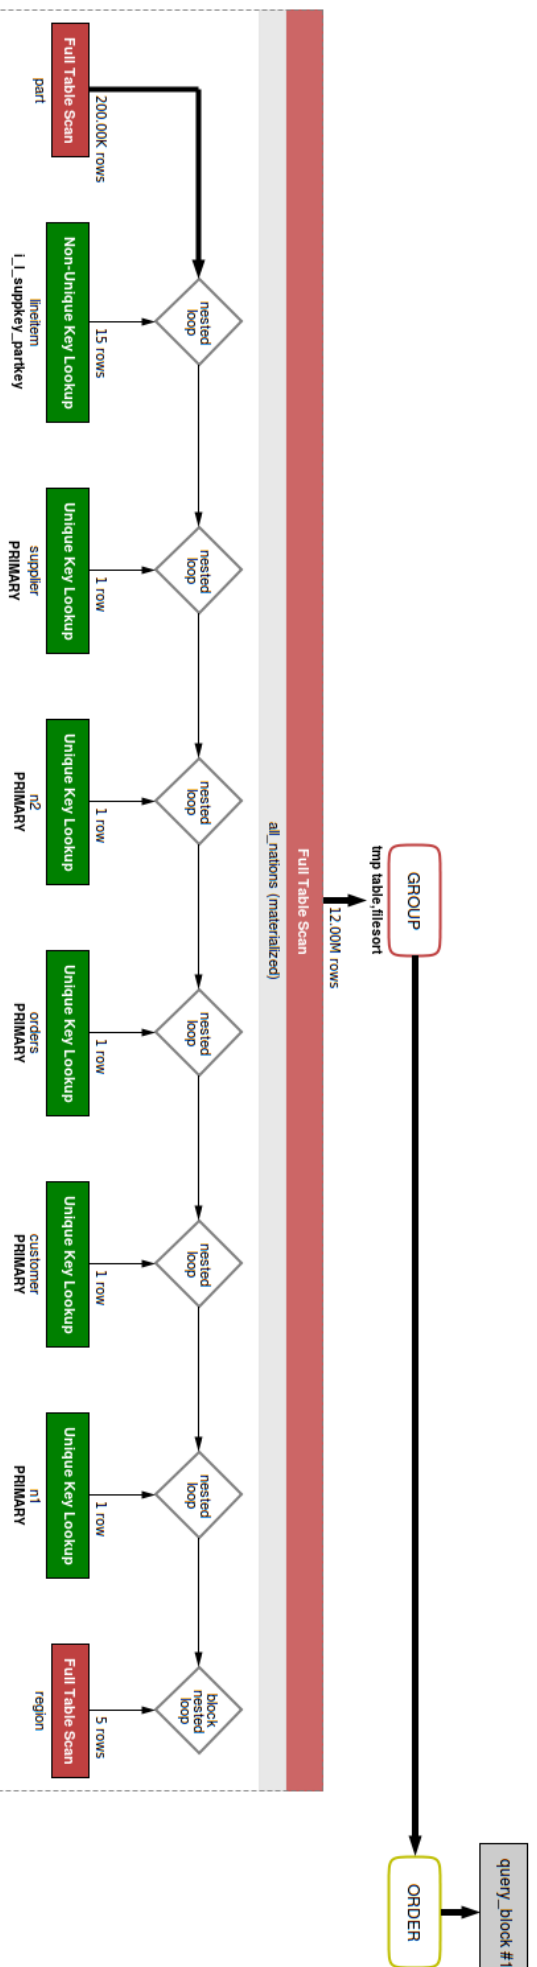
part before lineitem

Highest selectivity

# DBT-3 Query 8

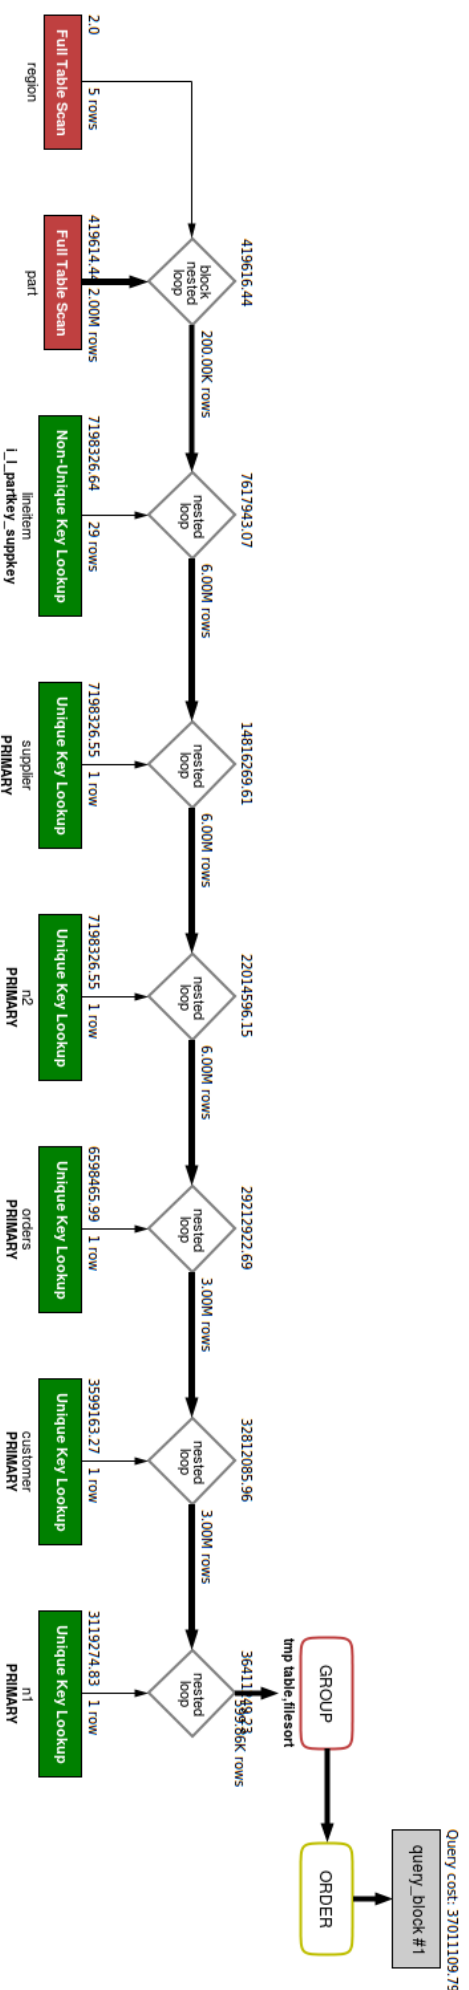
## Improved join order

Execution time: 3 seconds





# MySQL 5.7: Improved join order



## Improvements to Query 8 in MySQL 5.7:

- Filtering on non-indexed columns are taken into account
  - No need for hint to force **part** table to be processed early
- Merge derived tables into outer query
  - No temporary table

ORACLE®



# Program Agenda

- 1 Cost-based query optimization in MySQL
- 2 Tools for monitoring, analyzing, and tuning queries
- 3 Data access and index selection
- 4 Join optimizer
- 5 Subqueries
- 6 Sorting
- 7 Influencing the optimizer




# Overview of Subquery Optimizations

## Subquery category:

- IN (SELECT ...)
- NOT IN (SELECT ...)
- FROM (SELECT ...)
- <CompOp> ALL/ANY (SELECT ..)
- EXISTS/other

## Strategy:



# Traditional Optimization of IN Subqueries

IN → EXISTS transformation

- Convert IN subquery to EXISTS subquery by “push-down” IN-equality to subquery:

```
SELECT title FROM film
WHERE film_id IN (SELECT film_id FROM actor WHERE name="Bullock")
```

```
SELECT title FROM film
WHERE EXISTS (SELECT 1 FROM actor
              WHERE name="Bullock" AND film.film_id = actor.film_id)
```



- Benefit: subquery will evaluate fewer records
- Note: Special handling if pushed down expressions can be NULL

## Semi-join

- Convert subquery to inner join, BUT
  - Need some way to remove duplicates
- Different strategies for duplicate removal:
  - **FirstMatch** (equivalent to IN→EXISTS execution)
  - **LooseScan** (index scan, skip duplicates)
  - Materialization: **MatLookup** (like subquery materialization), **MatScan** (materialized table is first in join order)
  - **Duplicate WeedOut** (insert result rows of semi-join query into temporary table with unique index; duplicate rows will be rejected. Any join order.)
- If duplicate removal is not necessary:
  - **Table pull-out**

# Semi-join

## Continued


- Main advantage:
  - Opens up for more optimal "join orders".
  - Example:

```
SELECT o_orderdate, o_totalprice FROM orders
WHERE o_orderkey IN
(SELECT l_orderkey FROM lineitem WHERE l_shipDate='1996-09-30');
```

Will process less rows if starting with **lineitem** instead of **orders**
- Restriction:
  - Cannot use semi-join if subquery contains union or aggregation



# MySQL 5.6: Semi-join: Example



```
SELECT o_totalprice  
FROM orders  
WHERE o_orderkey IN  
      (SELECT l_orderkey  
       FROM lineitem  
       WHERE l_shipdate =  
             '1996-09-30');
```

DBT-3, Scale 10 (23 GB)  
innodb\_buffer\_pool\_size= 24 GB  
(CPU-bound)

# MySQL 5.7: Hint Example: SEMIJOIN



- No hint, optimizer chooses semi-join algorithm LooseScan:

**EXPLAIN SELECT \* FROM t2 WHERE t2.a IN (SELECT a FROM t3);**

id	select type	table	type	possible keys	key	key len	ref	rows	Extra
1	SIMPLE	t3	index	a	a	4	NULL	3	Using where; LooseScan
1	SIMPLE	t2	ref	a	a	4	test.t3.a	1	Using index

- Disable semi-join with hint:

**EXPLAIN SELECT \* FROM t2 WHERE t2.a IN (SELECT /\*+ NO\_SEMIJOIN() \*/ a FROM t3);**

id	select type	table	type	possible keys	key	key len	ref	rows	Extra
1	PRIMARY	t2	index	null	a	4	NULL	4	Using where; Using index
2	DEPENDENT SUBQUERY	t3	Index_subquery	a	a	4	func	1	Using index



# MySQL 5.7: Hint Example: SEMIJOIN

- Force Semi-join Materialization to be used

```
EXPLAIN SELECT /*+ SEMIJOIN(@subq MATERIALIZATION) */ * FROM t2
WHERE t2.a IN (SELECT /*+ QB_NAME(subq) */ a FROM t3);
```

id	select type	table	type	possible keys	key	key len	ref	rows	Extra
1	SIMPLE	t2	index	a	a	4	NULL	4	Using where; Using index
1	SIMPLE	<subquery2>	eq_ref	<auto_key>	<auto_key>	4	test.t2.a	1	NULL
2	MATERIALIZED	t3	index	a	a	4	NULL	3	Using index

# Subquery Materialization

1. Execute subquery once and store result in a temporary table
  - Table has unique index for quick look-up and duplicate removal.
2. Execute outer query and check for matches in temporary table.

```
SELECT o_orderdate, o_totalprice  
FROM orders  
WHERE o_orderkey IN (
```

```
    SELECT l_orderkey  
    FROM lineitem  
    GROUP BY l_orderkey  
    HAVING SUM(l_quantity) > 313
```

Materialize

```
);
```




# Comparing Subquery Materialization and IN → EXISTS

Query Execution Time Relative to MySQL 5.5

DBT-3, Scale 10 (23 GB)

innodb\_buffer\_pool\_size= 24 GB  
(CPU-bound)



ORACLE®

Copyright © 2016, Oracle and/or its affiliates. All rights reserved. |

# Subquery Materialization

```
SELECT o_orderdate, o_totalprice FROM orders WHERE o_orderkey IN (
SELECT l_orderkey FROM lineitem GROUP BY l_orderkey HAVING SUM(l_quantity) > 313);
```

id	select type	table	type	possible keys	key	key len	ref	rows	Extra
1	PRIMARY	orders	ALL	NULL	NULL	NULL	NULL	1500000	Using where
2	SUBQUERY	lineitem	index	PRIMARY, ...	PRIMARY	8	NULL	6001215	NULL

```
SELECT o_orderdate, o_totalprice FROM orders WHERE o_orderkey IN (SELECT /*+ SUBQUERY(INTOEXISTS)*/
l_orderkey FROM lineitem GROUP BY l_orderkey HAVING SUM(l_quantity) > 313);
```

id	select type	table	type	possible keys	key	key len	ref	rows	Extra
1	PRIMARY	orders	ALL	NULL	NULL	NULL	NULL	1500000	Using where
2	DEPENDENT SUBQUERY	lineitem	index	PRIMARY, ...	PRIMARY	8	NULL	6001215	NULL

## Derived Tables

- Subquery in FROM clause

```
SELECT AVG(o_totalprice) FROM  
( SELECT * FROM orders ORDER BY o_totalprice DESC LIMIT 100000 ) td;
```

- MySQL 5.6 and earlier: Executed separately and result stored in a temporary table (materialization)
- MySQL 5.7: Treat derived tables like views: May be merged with outer query block

# Index on Materialized Derived Table

Added in MySQL 5.6

```
SELECT o_clerk, price - o_totalprice  
FROM
```

```
(SELECT l_orderkey, SUM(l_extendedprice * (1 - l_discount)) price  
FROM lineitem GROUP BY l_orderkey) t1
```

```
JOIN
```

```
(SELECT o_clerk, o_orderkey, o_totalprice  
FROM orders WHERE o_orderdate BETWEEN '1995-01-01' AND '1995-12-31') t2  
ON t1.l_orderkey = t2.o_orderkey WHERE t1.price > t2.o_totalprice;
```

DBT-3 Scale Factor 10:

Create index for join

- MySQL 5.5: ? months; MySQL 5.6: 2 minutes

# Materialization of Derived Tables

## EXPLAIN

```
mysql> explain select o_clerk, price - o_totalprice from
(select l_orderkey, sum( l_extendedprice * (1 - l_discount)) price
  from lineitem group by l_orderkey) t1 join
(select o_clerk, o_orderkey, o_totalprice from orders
  where o_orderdate between '1995-01-01' and '1995-12-31') t2
  on t1.l_orderkey = t2.o_orderkey where t1.price > t2.o_totalprice;
```

id	select_type	table	type	possible_keys	key	
1	PRIMARY	<derived3>	ALL	NULL	NULL	...
1	PRIMARY	<derived2>	ref	<auto_key0>	<auto_key0>	...
3	DERIVED	orders	ALL	i_o_orderdate	NULL	...
2	DERIVED	lineitem	index	PRIMARY, ...	PRIMARY	...



## MySQL 5.7: Merge Derived Table with Outer Query

```
mysql> explain select o_clerk, price - o_totalprice from
      (select l_orderkey, sum( l_extendedprice * (1 - l_discount)) price
        from lineitem group by l_orderkey) t1 join
      (select o_clerk, o_orderkey, o_totalprice from orders
       where o_orderdate between '1995-01-01' and '1995-12-31') t2
 on t1.l_orderkey = t2.o_orderkey where t1.price > t2.o_totalprice;
```

id	select_type	table	type	possible_keys	key	
1	PRIMARY	<b>&lt;derived2&gt;</b>	ALL	NULL	NULL	...
1	PRIMARY	orders	eq_ref	PRIMARY, ...	PRIMARY	...
2	<b>DERIVED</b>	lineitem	index	PRIMARY, ...	PRIMARY	...

- Derived tables based on GROUP BY, DISTINCT, LIMIT, or aggregate functions will not be merged.
- MySQL 5.7: 1.5 minutes (DBT-3 SF10)

## Hint: Merge/Materialize Derived Table or View

MySQL 8.0.0 optimizer labs release

- Derived tables/views are, if possible, merged into outer query
- NO\_MERGE hint can be used to override default behavior:  

```
SELECT /*+ NO_MERGE(dt) */ *  
FROM t1 JOIN (SELECT x, y FROM t2) dt ON t1.x = dt.x;
```
- MERGE hint will force a merge  

```
SELECT /*+ MERGE(dt) */ *  
FROM t1 JOIN (SELECT x, y FROM t2) dt ON t1.x = dt.x;
```
- Can also use MERGE/NO\_MERGE hints for views  

```
SELECT /*+ NO_MERGE(v) */ * FROM t1 JOIN v ON t1.x = v.x;
```



# Program Agenda

- 1 Cost-based query optimization in MySQL
- 2 Tools for monitoring, analyzing, and tuning queries
- 3 Data access and index selection
- 4 Join optimizer
- 5 Subqueries
- 6 **Sorting**
- 7 Influencing the optimizer



# ORDER BY Optimizations

- General solution; “Filesort”:
  - Store query result in temporary table before sorting
  - If data volume is large, may need to sort in several passes with intermediate storage on disk.
- Optimizations:
  - Take advantage of index to generate query result in sorted order
  - For “LIMIT  $n$ ” queries, maintain priority queue of  $n$  top items in memory instead of filesort. (MySQL 5.6)



# Filesort

**SELECT \* FROM orders ORDER BY o\_totalprice ;**

id	select type	table	type	possible keys	key	key len	ref	rows	Extra
1	SIMPLE	orders	ALL	NULL	NULL	NULL	NULL	15000000	Using filesort

**SELECT c\_name, o\_orderkey, o\_totalprice  
FROM orders JOIN customer ON c\_custkey = o\_custkey  
WHERE c\_acctbal < -1000 ORDER BY o\_totalprice ;**

id	select type	table	type	possible keys	key	key len	ref	rows	Extra
1	SIMPLE	customer	ALL	PRIMARY	NULL	NULL	NULL	1500000	Using where; Using temporary; Using filesort
1	SIMPLE	orders	ref	i_o_custkey	i_o_custkey	5	...	7	NULL

# Filesort

## Status variables

Status variables related to sorting:

```
mysql> show status like 'Sort%';
```

Variable_name	Value
Sort_merge_passes	1
Sort_range	0
Sort_rows	136170
Sort_scan	1

>0: Intermediate storage on disk.

Consider increasing sort\_buffer\_size

Number of sort operations  
(range scan or table/index scans)

Number of rows sorted



# Filesort

## Performance Schema

Sorting status per statement available from Performance Schema

```
mysql> SELECT sql_text, sort_merge_passes, sort_range, sort_rows, sort_scan  
FROM performance_schema.events_statements_history  
ORDER BY timer_start DESC LIMIT 1;
```

sql_text	sort_merge_passes	sort_range	sort_rows	sort_scan
SELECT ...	1	0	136170	1

# Filesort: Case Study

```
mysql> FLUSH STATUS;

Query OK, 0 rows affected (0.00 sec)

mysql> SELECT AVG(o_totalprice) FROM
  ( SELECT * FROM orders
    ORDER BY o_totalprice DESC
    LIMIT 100000) td;
+-----+
| AVG(o_totalprice) |
+-----+
| 398185.986158     |
+-----+
1 row in set (24.65 sec)
```

Unnecessary large data volume!

```
mysql> SHOW STATUS LIKE 'sort%';
+-----+-----+
| Variable_name | Value |
+-----+-----+
| Sort_merge_passes | 1432 |
| Sort_range       | 0     |
| Sort_rows        | 100000 |
| Sort_scan        | 1     |
+-----+-----+
4 rows in set (0.00 sec)
```

Many intermediate sorting steps!



# Filesort: Case Study

Reduce amount of data to be sorted

```
mysql> SELECT AVG(o_totalprice) FROM (SELECT o_totalprice FROM orders ORDER BY
      o_totalprice DESC LIMIT 100000) td;
+-----+
| AVG(o_totalprice) |
+-----+
| 398185.986158      |
+-----+
1 row in set (8.18 sec)
```

```
mysql> SELECT sql_text, sort_merge_passes FROM performance_schema.
      events_statements_history ORDER BY timer_start DESC LIMIT 1;
```

```
+-----+-----+
| sql_text                                     | sort_merge_passes |
+-----+-----+
| SELECT AVG(o_totalprice) FROM (SELECT o_totalprice | 229 |
+-----+-----+
```



# Filesort: Case Study

Default is 256 kB

Increase sort buffer (1 MB)

```
mysql> SET sort_buffer_size = 1024*1024;
```

```
mysql> SELECT AVG(o_totalprice) FROM (SELECT o_totalprice FROM orders ORDER BY  
o_totalprice DESC LIMIT 100000) td;
```

```
+-----+  
| AVG(o_totalprice) |  
+-----+  
| 398185.986158      |  
+-----+
```

```
1 row in set (7.24 sec)
```

```
mysql> SELECT sql_text, sort_merge_passes FROM performance_schema.  
events_statements_history ORDER BY timer_start DESC LIMIT 1;
```

```
+-----+-----+  
| sql_text                                     | sort_merge_passes |  
+-----+-----+  
| SELECT AVG(o_totalprice) FROM (SELECT o_totalprice | 57 |  
+-----+-----+
```



# Filesort: Case Study

Increase sort buffer even more (8 MB)

```
mysql> SET sort_buffer_size = 8*1024*1024;
```

```
mysql> SELECT AVG(o_totalprice) FROM (SELECT o_totalprice FROM orders ORDER BY  
o_totalprice DESC LIMIT 100000) td;
```

```
+-----+  
| AVG(o_totalprice) |  
+-----+  
| 398185.986158      |  
+-----+
```

```
1 row in set (6.30 sec)
```

```
mysql> SELECT sql_text, sort_merge_passes FROM performance_schema.  
events_statements_history ORDER BY timer_start DESC LIMIT 1;
```

```
+-----+-----+  
| sql_text                                     | sort_merge_passes |  
+-----+-----+  
| SELECT AVG(o_totalprice) FROM (SELECT o_totalprice | 0 |  
+-----+-----+
```



# Using Index to Avoid Sorting

CREATE INDEX i\_o\_totalprice ON orders(o\_totalprice);

SELECT o\_orderkey, o\_totalprice FROM orders ORDER BY o\_totalprice ;

id	select type	table	type	possible keys	key	key len	ref	rows	Extra
1	SIMPLE	orders	index	NULL	i_o_totalprice	6	NULL	15000000	Using index

However, still (due to total cost):

SELECT \* FROM orders ORDER BY o\_totalprice ;

id	select type	table	type	possible keys	key	key len	ref	rows	Extra
1	SIMPLE	orders	ALL	NULL	NULL	NULL	NULL	15000000	Using filesort



# Using Index to Avoid Sorting

Case study revisited

```
SELECT AVG(o_totalprice) FROM  
(SELECT o_totalprice FROM orders ORDER BY o_totalprice DESC LIMIT 100000) td;
```

id	select type	table	Type	possible keys	key	key len	ref	rows	Extra
1	PRIMARY	<derived2>	ALL	NULL	NULL	NULL	NULL	100000	NULL
2	DERIVED	orders	index	NULL	i_o_totalprice	6	NULL	15000000	Using index

```
mysql> SELECT AVG(o_totalprice) FROM (  
  SELECT o_totalprice FROM orders  
  ORDER BY o_totalprice DESC LIMIT 100000) td;  
...  
1 row in set (0.06 sec)
```

ORACLE®

# Program Agenda

- 1 Cost-based query optimization in MySQL
- 2 Tools for monitoring, analyzing, and tuning queries
- 3 Data access and index selection
- 4 Join optimizer
- 5 Sorting
- 6 Influencing the optimizer



# Program Agenda

- 1 Cost-based query optimization in MySQL
- 2 Tools for monitoring, analyzing, and tuning queries
- 3 Data access and index selection
- 4 Join optimizer
- 5 Subqueries
- 6 Sorting
- 7 Influencing the optimizer

# Influencing the Optimizer

When the optimizer does not do what you want

- Add indexes
- Force use of specific indexes:
  - USE INDEX, FORCE INDEX, IGNORE INDEX
- Force specific join order:
  - STRAIGHT\_JOIN
- Adjust session variables
  - optimizer\_switch flags: set optimizer\_switch="index\_merge=off"
  - Buffer sizes: set sort\_buffer=8\*1024\*1024;
  - Other variables: set optimizer\_search\_depth = 10;

# MySQL 5.7: New Optimizer Hints

- New hint syntax:

```
SELECT /*+ HINT1(args) HINT2(args) */ ... FROM ...
```

- New hints:

- BKA(*tables*)/NO\_BKA(*tables*), BNL(*tables*)/NO\_BNL(*tables*)
- MRR(*table indexes*)/NO\_MRR(*table indexes*)
- SEMIJOIN/NO\_SEMIJOIN(*strategies*), SUBQUERY(*strategy*)
- NO\_ICP(*table indexes*)
- NO\_RANGE\_OPTIMIZATION(*table indexes*)
- QB\_NAME(*name*)

- Finer granularity than **optimizer\_switch** session variable

# Optimizer Hints

## Future

- New hints in 8.0.0 Optimizer Labs Release
  - Enable/disable merge of views and derived tables:
    - `MERGE() NO_MERGE()`
  - Join order
    - `JOIN_ORDER(tables) JOIN_PREFIX(tables) JOIN_SUFFIX(tables) JOIN_FIXED_ORDER()`
- Hints we consider to add
  - Force/ignore index\_merge alternatives
  - Reimplement index hints in new syntax
  - Temporarily set session variables for just one query

## MySQL 5.7: Query Rewrite Plugin

- Rewrite problematic queries without the need to make application changes
  - Add hints
  - Modify join order
  - Much more ...
- Add rewrite rules to table:

```
INSERT INTO query_rewrite.rewrite_rules (pattern, replacement ) VALUES
("SELECT * FROM t1 WHERE a > ? AND b = ?",
"SELECT * FROM t1 FORCE INDEX (a_idx) WHERE a > ? AND b = ?");
```
- New pre- and post-parse query rewrite APIs
  - Users can write their own plug-ins





# MySQL 5.7: Adjustable Cost Constants

Experimental! Use with caution! No guarantees!

```
EXPLAIN SELECT SUM(o_totalprice) FROM orders
WHERE o_orderdate BETWEEN '1994-01-01' AND '1994-12-31';
```

id	select type	table	type	possible keys	key	key len	rows	filtered	Extra
1	SIMPLE	orders	ALL	i_o_orderdate	NULL	NULL	15000000	29.93	Using where

```
UPDATE mysql.engine_cost SET cost_value=0.2 ← Default: 1.0
WHERE cost_name='memory_block_read_cost';
```

```
FLUSH COST_CONSTANTS; ← Make server read new cost constants
```



# MySQL 5.7: Adjustable Cost Constants

Continued

```
EXPLAIN SELECT SUM(o_totalprice) FROM orders  
WHERE o_orderdate BETWEEN '1994-01-01' AND '1994-12-31';
```

Id	select type	table	type	possible keys	key	key len	rows	filtered	Extra
1	SIMPLE	orders	range	i_o_orderdate	i_o_orderdate	4	4489990	100.00	Using index condition

Note:

- Heuristic used: If table is smaller than 20% of database buffer, all pages are in memory
- Only new connections will see updated cost constants

# More information

- MySQL Server Team blog
  - <http://mysqlserverteam.com/>
- My blog:
  - <http://oysteing.blogspot.com/>
- Optimizer team blog:
  - <http://mysqloptimizerteam.blogspot.com/>
- MySQL forums:
  - Optimizer & Parser: <http://forums.mysql.com/list.php?115>
  - Performance: <http://forums.mysql.com/list.php?24>



## Some Relevant Presentations at OpenWorld 2015

- MySQL Optimizer: What's New in 5.7 and Sneak Peek at 8.0 [CON6112]
  - Wednesday, Sep 21, 3:00 p.m. | Park Central - City
- MySQL 8.0: Common Table Expressions [CON7928]
  - Thursday, Sep 22, 12:00 p.m. | Park Central - Stanford
- A MySQL Sys Schema Deep Dive [CON4589]
  - Thursday, Sep 22, 1:15 p.m. | Park Central - City
- MySQL Performance Tuning 101 [CON6194]
  - Tuesday, Sep 20, 12:15 p.m. | Park Central - City
- Meet the MySQL Engineering Team [BOF1967]
  - Tuesday, Sep 20, 6:15 p.m. | Park Central - City

# General Session: Monday, 4:15pm, YBCA Theater

## State of the Dolphin



- Rich Mason, SVP & General Manager MySQL GBU, Oracle
- Tomas Ulin, VP MySQL Engineering, Oracle

## Customer Experiences

**Booking.com**



**CHURCHILL DOWNS**

Nicolai Plum, Senior Systems Architect, Booking.com

Andrew Archibald, VP of IT, Churchill Downs

**ORACLE**

# MySQL Community Reception @ Oracle OpenWorld

**Celebrate, Have Fun and Mingle with Oracle's MySQL Engineers & Your Peers**

- Tuesday, September 20 @ 7 pm
- Jillian's at Metreon: 175 Fourth Street, San Francisco  
At the corner of Howard and 4<sup>th</sup> st.; only 2-min walk from Moscone Center  
(same place as last year)

# Join us!


ORACLE®





# Oracle Support Stars Bar

- Ask the Experts your toughest product questions: MySQL & all Oracle products!
- View My Oracle Support and Product Demonstrations
- Learn what's new and more!
- Moscone West Exhibition Hall, Booth 3451



ORACLE®

[oracle.com/goto/starsbar](http://oracle.com/goto/starsbar)

Copyright © 2016, Oracle and/or its affiliates. All rights reserved.

## Safe Harbor Statement

The preceding is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, and timing of any features or functionality described for Oracle's products remains at the sole discretion of Oracle.



# Integrated Cloud

## Applications & Platform Services



ORACLE®